

HARSH PARIYA

// Agentic AI & Generative AI Engineer • Autonomous Multi-Agent Systems • LLMs

Hands-on **Agentic AI & Generative AI Engineer** with industry experience architecting **autonomous multi-agent execution graphs (LangGraph, LangChain)**, autonomous tool execution (**Model Context Protocol — MCP**), and production RAG pipelines with deterministic validation guardrails. Proven ability to build self-correcting agent systems, GraphRAG code intelligence, and full-stack AI platforms with sub-second retrieval. Active open-source engineer with **1,900+ GitHub contributions**.

Ahmedabad, Gujarat, India • +91 96019 86209
harshpariya195@gmail.com • harshpariya.dev • linkedin.com/in/harsh-pariya • github.com/HarshPariya

1,900+ GitHub contributions & open-source commits	10 Live KPIs LeadFlow-AI sales & CRM orchestration	0% Error 0.65 cosine similarity relevance gate	45% Cut in RAG query latency via vector caching
---	--	--	---

01 WHAT I BUILD

Autonomous agent architectures, multi-agent graphs & production AI platforms

Multi-Agent Orchestration LangGraph state machines, multi-turn reasoning loops, dynamic planning, hypothesis generators, and critic agents.	Tool Execution & MCP Model Context Protocol (MCP) tool calling, autonomous agent execution, and human-in-the-loop safety rails.	Production RAG & GraphRAG Token-level semantic chunking (300–500 tokens), AST code intelligence, Sentence Transformers, and RRF reranking.
Vector Search & Memory MongoDB Atlas Vector Search, Pinecone, ChromaDB, Qdrant, Redis semantic query-hash caching, and indexing.	Agentic Full-Stack Deployment High-throughput FastAPI backends, Server-Sent Events (SSE) streaming, Docker containerization, and Next.js 14 UIs.	Guardrails & Reliability Calibrated 0.65 cosine similarity gates eliminating hallucinations, Pydantic schemas, and context window optimization.

02 PROFESSIONAL EXPERIENCE

Agentic AI Intern • Codage Habitation, Ahmedabad (On-site)

Jul 2026 – Present

- Architecting and deploying enterprise **Agentic AI systems** using Python, LangGraph, and open-source foundation LLMs to automate multi-step operational reasoning.
- Engineering **state machine execution graphs** with dynamic runtime nodes (retrieve_only, tool_only, answer_directly), MCP tool bindings, and human-in-the-loop approval rails.
- Designing high-precision **RAG pipelines** with token-based chunking (300–500 tokens, 15% overlap) and vector similarity caching, cutting latency by **45%** and completely eliminating factual hallucinations.
- Conducting prompt & context engineering experiments; optimizing system instructions and context windows to maximize accuracy, determinism, and steerability across production LLM workflows.

Artificial Intelligence Intern • Codec Technologies India, Ahmedabad (Remote)

May 2026 – Jul 2026

- Developed and evaluated machine learning models and NLP data preprocessing pipelines using **Python, Scikit-Learn, Pandas, and NumPy**.
- Trained and benchmarked classification algorithms (**XGBoost, SVM, Random Forest**), tuning hyperparameters via cross-validation to achieve **90%+ baseline validation accuracy**.
- Containerized ML inference models using **FastAPI and Docker**, establishing REST endpoints for real-time inference with automated input schema validation.

03 EDUCATION

B.Tech, Computer Science & Engineering Rai University, Ahmedabad · CGPA: 8.98 / 10 · Core Coursework: Data Structures & Algorithms, OOP, Machine Learning, Deep Learning, Probability & Statistics, DBMS, Operating Systems, Software Engineering. Industrial Exposure: Adani Power Plant & Mundra Port (PLCs, SCADA, Control Systems & Perimeter Cybersecurity). Languages: English (Fluent), Hindi (Fluent), Gujarati (Native).	2024 – 2028 (Expected)
---	-------------------------------

04 DEPLOYED AI APPLICATIONS

Production platforms with verified live web endpoints & source repositories

AI Career & Research Platform Live Demo ↗ Frontend Repo ↗ Backend Repo ↗ Full-stack Generative AI platform integrating 6 AI engines: ATS Resume Reviewer, RAG PDF Chatbot, Research Agent, Vision Q&A, Interview Assistant. LangChain, Llama 3.1 & Llama 4 Scout Vision , Next.js + FastAPI. Next.js · FastAPI · Groq · LLaMA 3.1 · MongoDB Vector Search · Tailwind	SquidAI — Coding Assistant Live Demo ↗ GitHub Repo ↗ Full-stack LLM assistant designed for real-time automated code generation, syntax/logic bug explanations, and interactive technical Q&A with dynamic context tracking. React · Node.js · Generative AI · REST API · Vercel Deployment	California House Price AI Live Demo ↗ GitHub Repo ↗ End-to-end regression pipeline on California housing dataset achieving 85.03% R² score . Microservice deployed via containerized FastAPI backend with Next.js frontend. Python · XGBoost · FastAPI · Next.js · Scikit-Learn · Vercel
---	---	--

LeadFlow AI — Intelligent CRM & Sales Automation Platform

Production

Live Demo ↗

GitHub Repo ↗

- Architected an enterprise B2B sales automation platform that autonomously captures incoming prospects, enforces **normalized email duplicate prevention** in MongoDB Atlas, and executes real-time AI lead qualification.
- Integrated **Groq Llama 3.3 70B** for structured JSON scoring (0–100 assessing budget, authority, need, urgency) paired with deterministic priority tiers (HIGH, MEDIUM, LOW); built native REST adapters for **Twenty CRM** to idempotently upsert Persons, Companies, and Opportunities with automated follow-up tasks.
- Engineered hardened **Zapier catch-hook webhooks** with cryptographically unique event IDs (replay attack protection) and 3–path conditional routing for automated Gmail outreach; delivered an operational dashboard with **10 production KPIs** and dual-view (Kanban/Table) 7–stage deal tracking.

Next.js 15 (App Router) · React 19 · TypeScript · Groq LLaMA 3.3 70B · MongoDB Atlas · Twenty CRM · Zapier Webhooks · Tailwind CSS v4

Agentflow-AI (“Ask, Retrieve, Act”) — Autonomous Enterprise AI Agent

Production

GitHub Repo ↗

- Engineered an autonomous enterprise operations copilot using **LangGraph state machines** with an “Ask, Retrieve, Act” runtime loop dynamically routing queries between `retrieve_only`, `tool_only`, or direct synthesis.
- Integrated **Model Context Protocol (MCP)** client for external tool execution, guarded by a human-in-the-loop confirmation safety checkpoint for state-modifying actions.
- Built a dynamic PDF ingestion pipeline with **PyMuPDF**, applying semantic token chunking, **0.65 cosine relevance gate** eliminating hallucinations, Redis query-hash caching, and MongoDB Atlas step logging with Next.js 14 streaming.

Python · LangGraph · FastAPI (SSE) · Redis · MongoDB Atlas · Next.js 14 · PyMuPDF · Model Context Protocol (MCP) · Groq/OpenAI

06 ML, COMPUTER VISION & NLP SYSTEMS

Deployed models with verified Live Demos & GitHub source repositories

Brain Tumor Detection — CNN

Live Demo ↗

GitHub ↗

TensorFlow/Keras CNN classifying MRI brain scans into 4 categories: Glioma, Meningioma, Pituitary, and No Tumor. Trained from scratch on medical imaging dataset.

Python · TensorFlow · Keras · CNN · Vercel Deployment

Resume Screening AI — NLP Classifier

Live Demo ↗

GitHub ↗

Multi-category resume classifier using XGBoost and TF-IDF vectorization. Achieved **88%+ accuracy** across 24 job categories; deployed end-to-end on Streamlit.

Python · XGBoost · TF-IDF · NLP · Streamlit Production

Fake News Detection — NLP

Live Demo ↗

GitHub ↗

Linear SVM trained on 44,000+ news articles using TF-IDF features, achieving **99.29% test accuracy** with real-time interactive verification.

Python · Linear SVM · TF-IDF · NLP · Streamlit

Image Caption Generator — CNN + LSTM

Live Demo ↗

GitHub ↗

Encoder-decoder architecture combining InceptionV3 for visual feature extraction with an LSTM decoder for natural-language image captioning.

Python · TensorFlow · CNN (InceptionV3) · LSTM · Streamlit

Open Source & Problem Solving

Verified

- **1,900+ GitHub contributions** across 15+ ML, GenAI, and full-stack repositories.
- **200+ DSA problems** solved across LeetCode & GeeksforGeeks.
- Active hackathon competitor: AMD Slingshot Ideathon & Google Developer Group (GDG) AI events.

Certifications & Professional Training

Credentials

AI Tools & ChatGPT Workshop · Be10x (Completed)

View on LinkedIn ↗

Cloud DevOps Roadmap 2026 · GUVI x HCL (Webinar Completed)

View on LinkedIn ↗

- **Generative AI with Large Language Models:** DeepLearning.AI & AWS
- **Machine Learning Specialization:** DeepLearning.AI

07 CORE TECHNICAL SKILLS

Validated competencies across modern AI, frameworks, and deployment tooling

Agentic AI & LLMs

LangGraph · LangChain · Model Context Protocol (MCP) · Prompt Engineering
Context Window Optimization · Hallucination Mitigation · LoRA / PEFT SFT · Hugging Face
Groq API · LLaMA 3.3

RAG & Vector Search

MongoDB Atlas Vector Search · Pinecone · ChromaDB · Qdrant · Milvus · GraphRAG
Hybrid Search · Reciprocal Rank Fusion (RRF) · Sentence Transformers

Machine Learning & DL

Python · Scikit-Learn · PyTorch · TensorFlow · Keras · XGBoost · Random Forest
CNN (InceptionV3) · LSTM · SVM · Feature Engineering · TF-IDF

Backend & Deployment

FastAPI · Docker · Redis · Next.js 14 · Node.js · Express · REST APIs · SSE Streaming
Vercel · Streamlit · Linux · Git / GitHub · CI/CD

Data & Tooling

Pandas · NumPy · PyMuPDF · AST Parsing · Pydantic · n8n Automation · Cursor
Claude Code · GitHub Copilot